

Project Report on

LOAN ELIGIBILITY PREDICTION

Submitted By
Kalyani Avhale - 242062001

Guided By
Prof. Ashwini Matange

Post-Graduation Diploma in
Data Science & Artificial Intelligence

CONTENT

1. Problem Statement

2. Data Dictionary

3. R-Libraries Used

4. Understanding Data

5. Pre-processing /Cleaning

6. Handling Missing Values

7. Exploratory Data Analysis

8. Categorical Features

9. Model Building

I. Train Test Split

II. Logistic Regression

III. Random Forest

IV. Gradient Boosting

V. Predictions

10. Summary

PROBLEM STATEMENT

Dream Housing Finance company deals in all home loans. They have a presence across all urban, semi-urban, and rural areas. Customer-first applies for a home loan after that company validates the customer eligibility for a loan.

The company wants to automate the loan eligibility process (real-time) based on customer detail provided while filling the online application form. These details are Gender, Marital Status, Education, Number of Dependents, Income, Loan Amount, Credit History, and others. To automate this process, they have given a problem to identify the customer's segments, those are eligible for loan amount so that they can specifically target these customers. Here they have provided a partial data set.

DATA DICTIONARY

| Variable | Description |
|-------------------|--|
| Loan_ID | Unique Loan ID |
| Gender | Male/ Female |
| Married | Applicant married (Y/N) |
| Dependents | Number of dependents |
| Education | Applicant Education (Graduate/ Under Graduate) |
| Self_Employed | Self employed (Y/N) |
| ApplicantIncome | Applicant income |
| CoapplicantIncome | Coapplicant income |
| LoanAmount | Loan amount in thousands |
| Loan_Amount_Term | Term of loan in months |
| Credit_History | credit history meets guidelines |
| Property_Area | Urban/ Semi Urban/ Rural |
| Loan_Status | (Target) Loan approved (Y/N) |

#Loan_Status : is target feature

R LIBRARIES USED

```
# @Script for Loan Eligibility Prediction Script
# @Author : Kalyani Avhale
# @Language : R
# @Dataset source : https://www.kaggle.com/vikasukani/loan-eligible-dataset
# @Date : 17th April , 2021
```

```
setwd("D:/Trisem_2/R/Project/packages")
```

```
#Install Required Packages
```

```
# install.packages('tidyverse') # metapackage of all tidyverse packages
```

```
# install.packages('dplyr') #data manipulation(included in tidyverse)
```

```
# install.packages('caret') #for Classification and regression
```

```
# install.packages('ggpubr') #arranging plots into grids
```

```
# install.packages('modeest') #Estimation of the mode
```

```
# install.packages('ggplot2') #for plotting graphs(included in tidyverse)
```

```
# install.packages('ggcorrplot')#for plotting correlation matrix
```

```
# install.packages('randomForest')#Random forest
```

```
# install.packages('xgboost') #Gradient Boosting
```

```
#import packages
```

- library('tidyverse')
 - library('caret')
 - library('ggpubr')
 - library('modeest')
 - library('ggcorrplot')
 - library('randomForest')
 - library('xgboost')
-
-
-

UNDERSTANDING THE DATA

- `data=read.csv('loan-train.csv',na.strings=c(""))` #Load the dataset
- `head(data)`

```
  Loan_ID Gender Married Dependents Education Self_Employed ApplicantIncome CoapplicantIncome LoanAmount Loan_Amount_Term Credit_History
1 LP001002  Male    No         0 Graduate          No           5849              0          NA             360             1
2 LP001003  Male    Yes         1 Graduate          No           4583             1508          128             360             1
3 LP001005  Male    Yes         0 Graduate          Yes           3000              0           66             360             1
4 LP001006  Male    Yes         0 Not Graduate        No           2583             2358          120             360             1
5 LP001008  Male    No         0 Graduate          No           6000              0          141             360             1
6 LP001011  Male    Yes         2 Graduate          Yes           5417             4196          267             360             1
Property_Area Loan_Status
1      Urban            Y
2      Rural            N
3      Urban            Y
4      Urban            Y
5      Urban            Y
6      Urban            Y
> |
```

- `dim(data)` #Dimension of dataset
[1] 614 13
- `str(data)` #returns type of attribute along with firstfew values

```
'data.frame': 614 obs. of 13 variables:
 $ Loan_ID      : chr  "LP001002" "LP001003" "LP001005" "LP001006" ...
 $ Gender       : chr  "Male" "Male" "Male" "Male" ...
 $ Married      : chr  "No" "Yes" "Yes" "Yes" ...
 $ Dependents   : chr  "0" "1" "0" "0" ...
 $ Education    : chr  "Graduate" "Graduate" "Graduate" "Not Graduate" ...
 $ Self_Employed : chr  "No" "No" "Yes" "No" ...
 $ ApplicantIncome : int  5849 4583 3000 2583 6000 5417 2333 3036 4006 12841 ...
 $ CoapplicantIncome: num  0 1508 0 2358 0 ...
 $ LoanAmount   : int  NA 128 66 120 141 267 95 158 168 349 ...
 $ Loan_Amount_Term : int  360 360 360 360 360 360 360 360 360 360 ...
 $ Credit_History : int  1 1 1 1 1 1 1 0 1 1 ...
 $ Property_Area : chr  "Urban" "Rural" "Urban" "Urban" ...
 $ Loan_Status  : chr  "Y" "N" "Y" "Y" ...
> |
```

- `attr_type <- sapply(data,class)` # list types for each attribute

```
Loan_ID      attr_type
Gender       character
Married      character
Dependents   character
Education    character
Self_Employed character
ApplicantIncome integer
CoapplicantIncome numeric
LoanAmount   integer
Loan_Amount_Term integer
Credit_History integer
Property_Area character
Loan_Status  character
> |
```

#Levels of classes

- `unique(data$Loan_Status)` # Since we have 2 classes it's a binary classification problem

```
[1] "Y" "N"
```

#class distribution

- `percent = prop.table(table(data$Loan_Status))*100`
- `cbind(freq=table(data$Loan_Status),percentage=percent)`

```
freq percentage
N 192 31.27036
Y 422 68.72964
```

#We see that only 31% of all the people in the dataset had a loan being approved. This means that our baseline model has an accuracy of 69%. An important measure when evaluating our model we be the sensitivity (aka recall aka the probability of detection as positive). If this value is low then our model is not very good at detecting true positive cases, even if the accuracy is very high. There are several ways to deal with imbalance

#Statistical Summary

- `summary(data)`

```
summary(data)
Loan_ID      Gender      Married      Dependents      Education      Self_Employed      ApplicantIncome
Length:614   Length:614   Length:614   Length:614   Length:614   Length:614   Min.   : 150
Class :character Class :character Class :character Class :character Class :character Class :character 1st Qu.: 2878
Mode  :character Mode  :character Mode  :character Mode  :character Mode  :character Mode  :character Median : 3812
                                         Mean   : 5403
                                         3rd Qu.: 5795
                                         Max.   :81000

CoapplicantIncome  LoanAmount  Loan_Amount_Term  Credit_History  Property_Area  Loan_Status
Min.   : 0         Min.   : 9.0      Min.   : 12       Min.   :0.0000   Length:614     Length:614
1st Qu.: 0         1st Qu.:100.0    1st Qu.:360      1st Qu.:1.0000   Class :character Class :character
Median : 1188      Median :128.0    Median :360      Median :1.0000   Mode  :character Mode  :character
Mean   : 1621      Mean   :146.4    Mean   :342      Mean   :0.8422
3rd Qu.: 2297      3rd Qu.:168.0    3rd Qu.:360      3rd Qu.:1.0000
Max.   :41667      Max.   :700.0    Max.   :480      Max.   :1.0000
NA's   :22         NA's   :14       NA's   :50
```

we can see few cols has NA's and the scale for variables differ

We can perform scaling in later steps

We can see few outlier values

PRE-PROCESSING/CLEANING

#check for duplicate rows

- `dim(loan_dataset[duplicated(loan_dataset$Loan_ID),])`
[1] 0 13

#we have 0 duplicates across 13 columns for Loan_Id

#check for all unique values across dataset (probably for char type values)

- `Col_names=c("Gender","Married","Dependents","Education",
"Self_Employed","Loan_Amount_Term",
"Credit_History","Property_Area","Loan_Status")`
- `lapply(loan_dataset[Col_names], function(x) unique(x))`

```
$Gender
[1] "Male" "Female" NA

$Married
[1] "No" "Yes" NA

$Dependents
[1] "0" "1" "2" "3+" NA

$Education
[1] "Graduate" "Not Graduate"

$Self_Employed
[1] "No" "Yes" NA

$Loan_Amount_Term
[1] 360 120 240 NA 180 60 300 480 36 84 12

$Credit_History
[1] 1 0 NA

$Property_Area
[1] "Urban" "Rural" "Semiurban"

$Loan_Status
[1] "Y" "N"
```

Dependents,Gender,Self_Employed has NA's and 3 suffixed with +

Loan_Amount_Term,Credit_History has NA's

#check for null values count for numerical type cols

- `colSums(is.na(loan_dataset))`

| Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome | CoapplicantIncome |
|------------|------------------|----------------|---------------|-------------|---------------|-----------------|-------------------|
| 0 | 13 | 3 | 15 | 0 | 32 | 0 | 0 |
| LoanAmount | Loan_Amount_Term | Credit_History | Property_Area | Loan_Status | | | |
| 22 | 14 | 50 | 0 | 0 | | | |

#replace 3+ with 3

- `loan_dataset$Dependents <-
replace(loan_dataset$Dependents,loan_dataset$Dependents=='3+',3)`

HANDLING MISSING VALUES

- `numeric_cols <-
c('ApplicantIncome','CoapplicantIncome','LoanAmount','Loan_Amount_Term')
#numerical cols`
- `cat_cols <- c('Gender','Married','Dependents','Self_Employed','Credit_History')
#categorical cols`

`#numeric NA : mean`

- `data_beforeImputation[numeric_cols] <-
apply(data_beforeImputation[numeric_cols], function(x)ifelse(is.na(x), mean(x,
na.rm=TRUE), x))`
- `colSums(is.na(data_beforeImputation))`

`#categorical NA : fill with mode`

- `data_beforeImputation[cat_cols] <- sapply(data_beforeImputation[cat_cols],
function(x)ifelse(is.na(x), mfv(x), x))`
- `colSums(is.na(data_beforeImputation))`

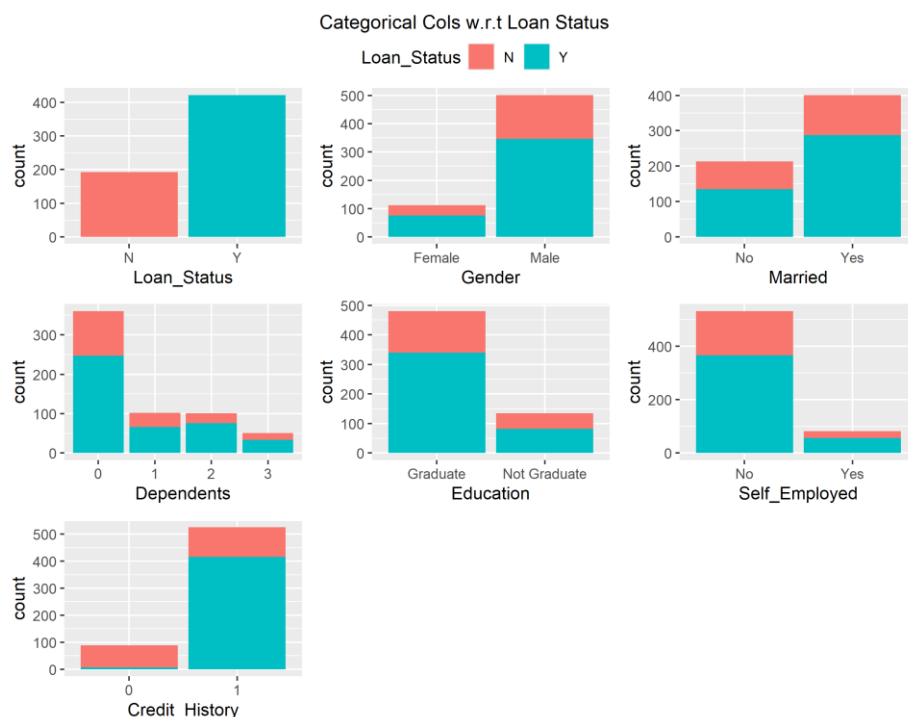
| Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome | CoapplicantIncome |
|------------|------------------|----------------|---------------|-------------|---------------|-----------------|-------------------|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| LoanAmount | Loan_Amount_Term | Credit_History | Property_Area | Loan_Status | | | |
| 0 | 0 | 0 | 0 | 0 | | | |

EXPLORATORY DATA ANALYSIS

➤ `setwd("D:/Trisem_2/R/Project/Plots")` #derictory to save plots

#visualizaing categorical variables first with respect to Loan Status

- `ls_plt <- ggplot(data = data_afterImputation) +
geom_bar(mapping = aes(x = Loan_Status,fill=Loan_Status))`
- `g_plt <- ggplot(data = data_afterImputation) +
geom_bar(mapping = aes(x = Gender,fill=Loan_Status))`
- `m_plt <- ggplot(data = data_afterImputation) +
geom_bar(mapping = aes(x = Married,fill=Loan_Status))`
- `d_plt <- ggplot(data = data_afterImputation) +
geom_bar(mapping = aes(x = Dependents,fill=Loan_Status))`
- `e_plt <- ggplot(data = data_afterImputation) +
geom_bar(mapping = aes(x = Education,fill=Loan_Status))`
- `se_plt <- ggplot(data = data_afterImputation) +
geom_bar(mapping = aes(x = Self_Employed,fill=Loan_Status))`
- `ch_plt <- ggplot(data = data_afterImputation,aes(x=Credit_History,fill=Loan_Status))
+geom_bar()`
- `ggarrange(ls_plt,g_plt,m_plt,d_plt,e_plt,se_plt,ch_plt,nrow=3,ncol=3,common.legend
=TRUE)`




Inference :Male applicant has high loan approval,Applicant with 0 dependents has been approved with loan as compared to applicants with dependents. Self-employed applicants with loan approval e low as compared with which are not self employed(can be with other profession such as jobs,business,etc). Applicant with Credit History has highest loan approval

#Analyzing the three continuous variables w.r.t Loan_Status:


#ApplicantIncome,CoapplicantIncome,LoanAmount,Loan_Amount_Term

- ai_plt <- ggplot(data_afterImputation, aes(y= ApplicantIncome, x = "", fill = Loan_Status)) + geom_boxplot()
- cai_plt <- ggplot(data_afterImputation, aes(y= CoapplicantIncome, x = "", fill = Loan_Status)) + geom_boxplot()
- la_plt <- ggplot(data_afterImputation, aes(y= LoanAmount, x = "", fill = Loan_Status)) + geom_boxplot()
- lat_plt <- ggplot(data_afterImputation, aes(y= Loan_Amount_Term, x = "", fill = Loan_Status)) + geom_boxplot()
- figure <-
ggarrange(ai_plt,cai_plt,la_plt,lat_plt,nrow=2,ncol=2,common.legend=TRUE)
- annotate_figure(figure,top = "Continuous Cols w.r.t Loan Status")
- ggsave('Numerical_col_plot.png')



#Pair Plot

- `numeric_cols <-`
`c('ApplicantIncome','CoapplicantIncome','LoanAmount','Loan_Amount_Term')`
- `options(repr.plot.width =10, repr.plot.height = 10)` #adjust size of plot
- `pairs(data_afterImputation[numeric_cols], pch =`
`19,cex.labels=1.5,lower.panel=NULL)`



#We have positive correlations: LoanAmount and ApplicantIncome and LoanAmount and CoapplicantIncome

- `options(repr.plot.width =10, repr.plot.height = 10)` #adjust size of plot
- `par(mfrow=c(2,2))` #arrange plot (matrix grid)
- `boxplot(data_afterImputation$LoanAmount, horizontal=TRUE, main="Loan Amount", col='red')`
- `plot(density(data_afterImputation$LoanAmount), main="Density Graph", col='red')`
- `boxplot(log(data_afterImputation$LoanAmount), horizontal=TRUE, main="Loan Amount after log", col='blue')`
- `plot(density(log(data_afterImputation$LoanAmount)), main="Density Graph after log", col='blue')`

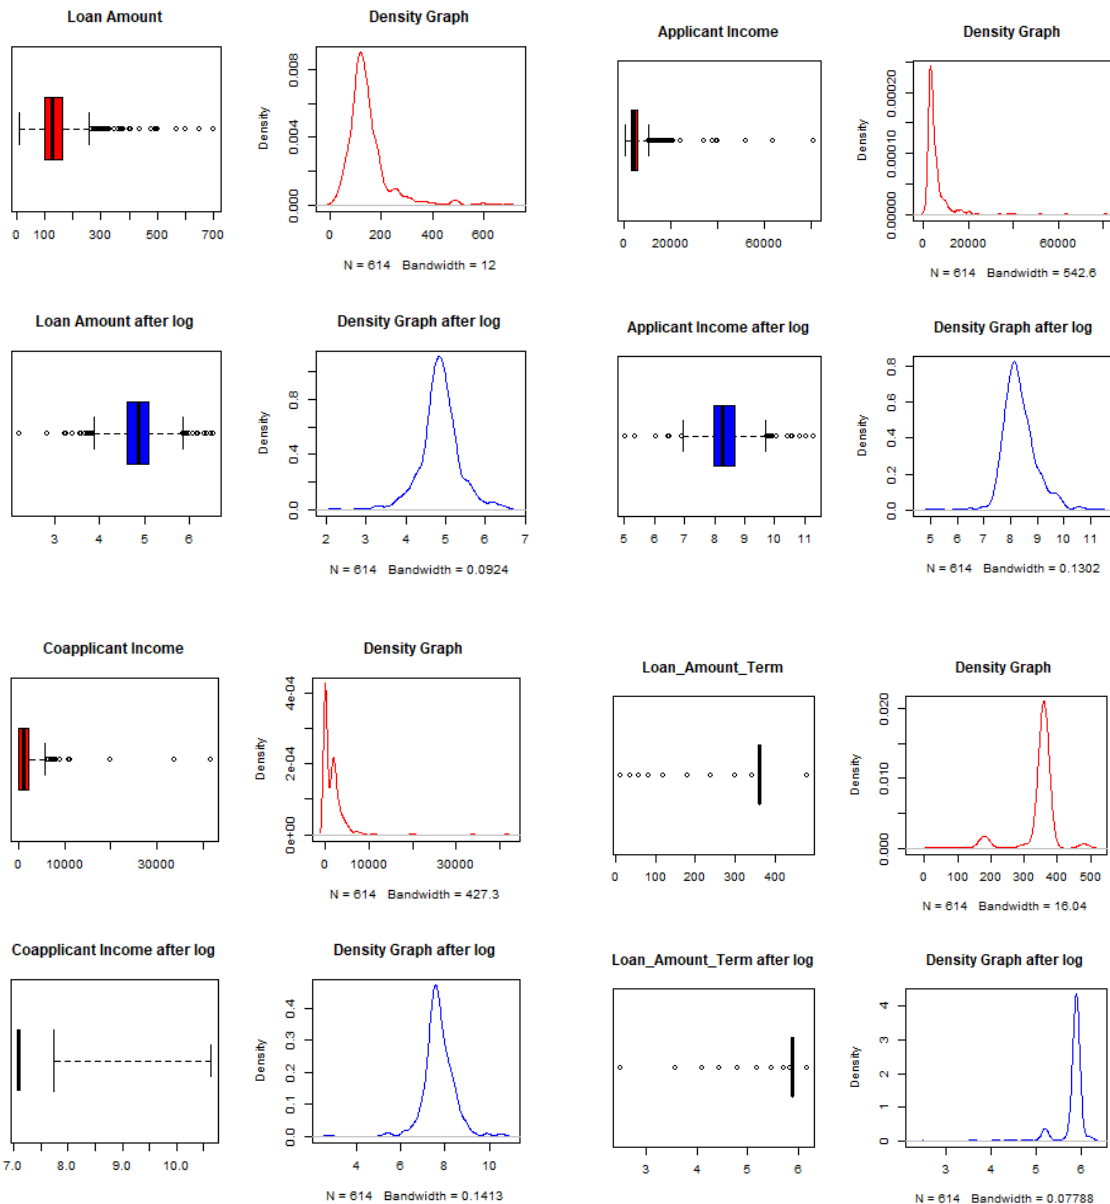
#LoanAmount log(removed the skewness)

#Similar plots for :

#ApplicantIncome

#CoapplicantIncome

#Loan_Amount_Term



#add log values to data

- `data_afterImputation$ApplicantIncome_log = log(data_afterImputation$ApplicantIncome)`
- `data_afterImputation$LoanAmount_log = log(data_afterImputation$LoanAmount)`
- `data_afterImputation$CoapplicantIncome_log = log(data_afterImputation$CoapplicantIncome)`

#log(0) is -Inf , so replace -Inf to 0

- `data_afterImputation[data_afterImputation== -Inf]<-0`

#lets check the covariance metrics

➤ `round(cor(data_afterImputation[numeric_cols]),3)`


we have +ve correlations : LoanAmount and ApplicantIncome and LoanAmount and coapplicantIncome

As Loan Amount depends on Income of applicant ,the more the Income has high probability of getting more Loan amount

| | ApplicantIncome | CoapplicantIncome | LoanAmount | Loan_Amount_Term |
|-------------------|-----------------|-------------------|------------|------------------|
| ApplicantIncome | 1.000 | -0.117 | 0.566 | -0.045 |
| CoapplicantIncome | -0.117 | 1.000 | 0.188 | -0.060 |
| LoanAmount | 0.566 | 0.188 | 1.000 | 0.039 |
| Loan_Amount_Term | -0.045 | -0.060 | 0.039 | 1.000 |

#co-relation Plot after scaling data

- `options(repr.plot.width =20, repr.plot.height = 20)#adjust size of plot`
- `loan_corr <- round(cor(af_scale),3) #get corr matrix`
- `ggcorrplot(loan_corr, title = "Correlation Plot",type = "lower",lab=TRUE,insig = "blank")`



Credit History has Positive correlation with target feature Loan Status

Dependents and Married , Loan Amount and Applicant Income are positively correlated

Applicant Income and CoApplicant Income has negative correlation

CATEGORICAL FEATURES

#drop log cols

- `loanData <- subset(data_afterImputation,select = -
c(ApplicantIncome,LoanAmount,CoapplicantIncome))`

#Education

- `loanData$Education <- sapply(loanData$Education,function(x)
ifelse(x=='Graduate',1,0)) #replace "Graduate" with 1 and "Not Graduate" with 0`

#Loan_status

- `loanData$Loan_Status <- sapply(loanData$Loan_Status,function(x)
ifelse(x=='Y',1,0))`

#Gender

- `loanData$Gender <- sapply(loanData$Gender,function(x) ifelse(x=='Male',1,0))
#replace "Y" with 1 and "N" with 0`

#Married

- `loanData$Married <- sapply(loanData$Married,function(x) ifelse(x=='Yes',1,0))
#replace "Y" with 1 and "N" with 0`

#Self_Employed

- `loanData$Self_Employed <- sapply(loanData$Self_Employed,function(x)
ifelse(x=='Yes',1,0)) #replace "Yes" with 1 and "No" with 0`

#Property_Area

- `loanData$Property_Area <- as.integer(factor(loanData$Property_Area))`

#Credit History and Dependents convert to numeric

- `loanData$Credit_History<-as.integer(loanData$Credit_History)`
- `loanData$Dependents<-as.integer(loanData$Dependents)`
- `cleanData_bfscale <- subset(loanData,select = -c(Loan_ID,Loan_Status))`

#Scaling

- `af_scale <- data.frame(scale(cleanData_bfscale))`
- `af_scale$Loan_Status <-loanData$Loan_Status`

MODEL BUILDING

1. Train Test Split

#Splitting the dataset into the Training set and Test set

- `set.seed(100) #randomization``
- `train_sample <- sample(nrow(af_scale), 0.75 * nrow(af_scale))`

#splitting data into training/testing data using the trainIndex object

- `trainData <- af_scale[train_sample,] #training data (75% of data)`
- `testData <- af_scale[-train_sample,] #testing data (25% of data)`

Check whether data set fairly even split

- `prop.table(table(trainData$Loan_Status))#train`
- `prop.table(table(testData$Loan_Status))#test`

```
> # Check whether data set fairly even split
> prop.table(table(trainData$Loan_Status))#train
      0      1
0.3108696 0.6891304
> prop.table(table(testData$Loan_Status))#test
      0      1
0.3181818 0.6818182
> |
```


2. LOGISTIC REGRESSION

glm() --> for generalized linear model and can be used to compute Logistic Regression

family = binomial is specified to perform binary classification

Predictions can be easily made using the function predict(). Use the option type = "response" to directly obtain the probabilities

glm model with all Features

- model_all <- glm(Loan_Status ~., data = trainData, family = binomial)
- pred_all <- predict.glm(model_all, testData[-12], type = 'response')
- pclass_all <- ifelse(pred_all < 0.5, 0, 1)
- confusionMatrix(table(as.factor(testData\$Loan_Status), pclass_all), positive = '1')

```
Confusion Matrix and Statistics

 pclass_all
 0      1
0   18   31
1     1  104

      Accuracy : 0.7922
    95% CI : (0.7195, 0.8533)
 No Information Rate : 0.8766
P-Value [Acc > NIR] : 0.9989

      Kappa : 0.4276

McNemar's Test P-Value : 2.951e-07

      Sensitivity : 0.7704
      Specificity : 0.9474
   Pos Pred Value : 0.9905
   Neg Pred Value : 0.3673
      Prevalence : 0.8766
Detection Rate : 0.6753
Detection Prevalence : 0.6818
Balanced Accuracy : 0.8589

'Positive' Class : 1
```

Features with minimum p value are good features

glm model with single feature Credit_History

- model1ch <- glm(Loan_Status ~Credit_History, data = trainData, family = binomial)
- pred_m1 <- predict.glm(model1ch, testData[-12], type = 'response')
- pclass_m1 <- ifelse(pred_m1 < 0.5, 0, 1)
- confusionMatrix(table(as.factor(testData\$Loan_Status), pclass_m1), positive = '1')

```
Confusion Matrix and Statistics

 pclass_m1
 0      1
0   18   31
1     1  104

      Accuracy : 0.7922
    95% CI : (0.7195, 0.8533)
 No Information Rate : 0.8766
P-Value [Acc > NIR] : 0.9989

      Kappa : 0.4276

McNemar's Test P-Value : 2.951e-07

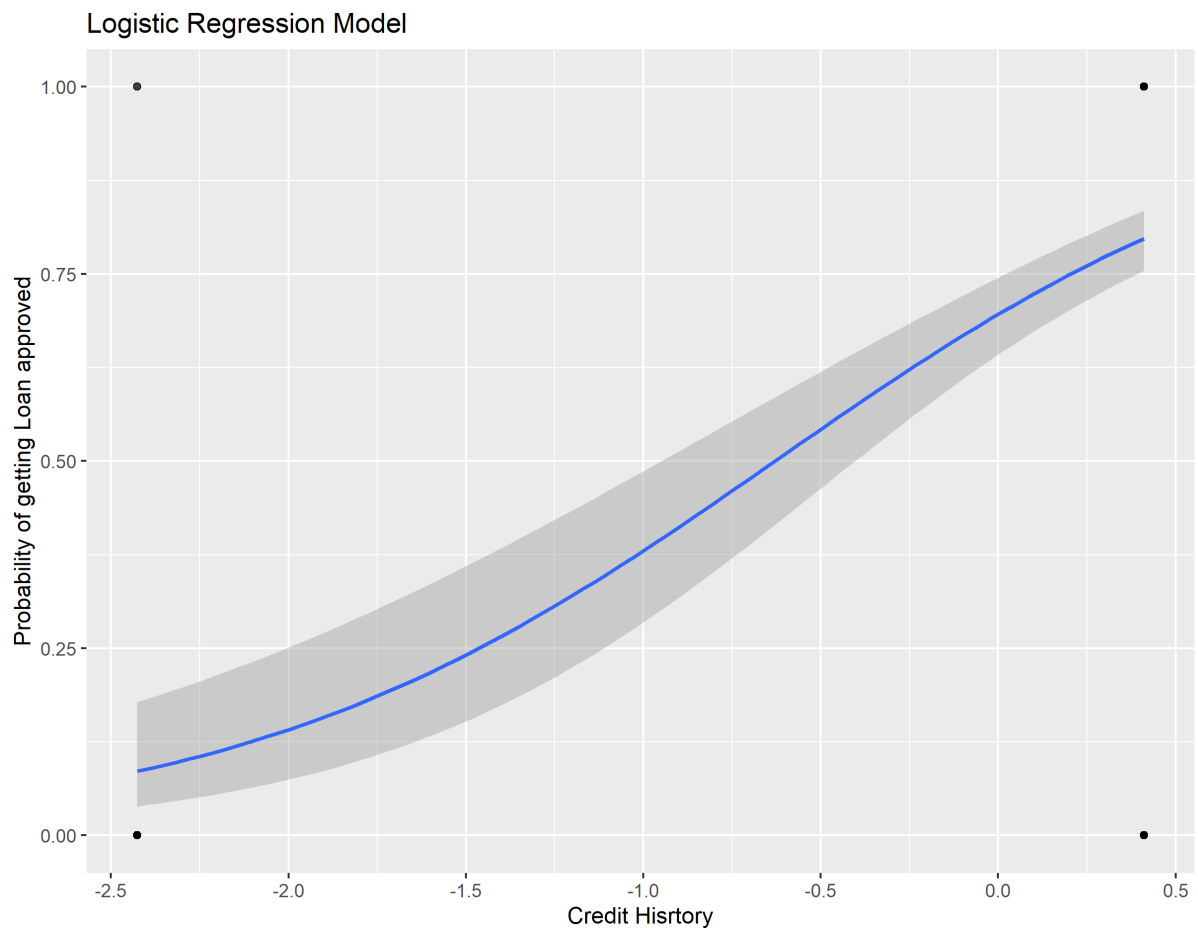
      Sensitivity : 0.7704
      Specificity : 0.9474
   Pos Pred Value : 0.9905
   Neg Pred Value : 0.3673
      Prevalence : 0.8766
Detection Rate : 0.6753
Detection Prevalence : 0.6818
Balanced Accuracy : 0.8589

'Positive' Class : 1
```

#Logistic curve

```
➤ ggplot(data=trainData,aes(Credit_History, Loan_Status)) +  
  geom_point(alpha = 0.2) +  
  geom_smooth(method = "glm", method.args = list(family = "binomial")) +  
  labs(  
    title = "Logistic Regression Model",  
    x = "Credit Histrory",  
    y = "Probability of getting Loan approved"  
  )  
➤ ggsave('Logistic_curve.png')
```

#No improvement in model accuracy



3. Random Forest

Bagging stands for bootstrap aggregating. It consists of building multiple different decision tree models from a single training data set by repeatedly using multiple bootstrapped subsets of the data and averaging the models. Here, each tree is build independently to the others.

Random Forest algorithm, is one of the most commonly used and the most powerful machine learning techniques. It is a special type of bagging applied to decision trees.

#random forest with default parameters and all features

- set.seed(100)
- original_rf<-randomForest(as.factor(Loan_Status)~ ., trainData,OOB=TRUE)
- original_rf

```
Call:
randomForest(formula = as.factor(Loan_Status) ~ ., data = trainData,      OOB = TRUE)
  Type of random forest: Classification
    Number of trees: 500
No. of variables tried at each split: 3

    OOB estimate of  error rate: 20.22%
Confusion matrix:
  0  1 class.error
0 66  77  0.53846154
1 16 301  0.05047319
> |
```

- pred <- predict(original_rf,testData[-12])
- confusionMatrix(as.factor(testData\$Loan_Status),pred)

```
Confusion Matrix and Statistics

      Reference
Prediction  0   1
  0   19  30
  1    2 103

      Accuracy : 0.7922
      95% CI   : (0.7195, 0.8533)
  No Information Rate : 0.8636
  P-Value [Acc > NIR] : 0.9947

      Kappa : 0.435

  McNemar's Test P-Value : 1.815e-06

      Sensitivity : 0.9048
      Specificity : 0.7744
   Pos Pred Value : 0.3878
   Neg Pred Value : 0.9810
      Prevalence : 0.1364
  Detection Rate : 0.1234
  Detection Prevalence : 0.3182
   Balanced Accuracy : 0.8396

  'Positive' Class : 0
```

- accuracy = sum(testData\$Loan_Status == pred)/length(testData\$Loan_Status)
#0.7922078
#no improvement observed as compared to glm model

Hyper Parameter Tunning – Random Forest

Hyperparameter tuning for Random Forest

```
➤ set.seed(10)
➤ tune_grid<-expand.grid(mtry=c(1:10), ntree=c(500,1000,1500,2000,2500,3000))
  #expand a grid of parameters
➤ mtry<-tune_grid[[1]]
➤ ntree<-tune_grid[[2]] #using vectors instead of dataframe to subset is faster in for
  loop
➤ OOB<-NULL #use to store calculated OOB error estimate

➤ for(i in 1:nrow(tune_grid)){
  rf<-randomForest(as.factor(Loan_Status)~. ,trainData, mtry=mtry[i], ntree=ntree[i])
  confusion<-rf$confusion
  temp<-(confusion[2]+confusion[3])/614 #calculate the OOB error estimate
  OOB<-append(OOB,temp)
}
```

```
➤ tune_grid$OOB<-OOB
➤ head(tune_grid[order(tune_grid["OOB"]), ], 4) #order the results
```

```
#   mtry ntree   OOB
# 22    2 1500 0.1400651
# 52    2 3000 0.1400651
#  2    2   500 0.1416938
# 32    2 2000 0.1416938
```

#We have optimal paramater with lowest OOB score mtry:2 and ntree:1500

#fit model with optimal Parameters

```
➤ final_rf <- randomForest(as.factor(Loan_Status)~. ,trainData, mtry=2, ntree=1500)
```

```
> final_rf
Call:
randomForest(formula = as.factor(Loan_Status) ~ ., data = trainData,      mtry = 2, ntree = 1500)
      Type of random forest: classification
      Number of trees: 1500
No. of variables tried at each split: 2

      OOB estimate of  error rate: 18.91%
Confusion matrix:
      0    1 class.error
0 64   79  0.55244755
1  8 309  0.02523659
> |
```

#Comparing the above random forest models :

#model original_rf has better train accuracy but low test accuracy (its getting over fitted)

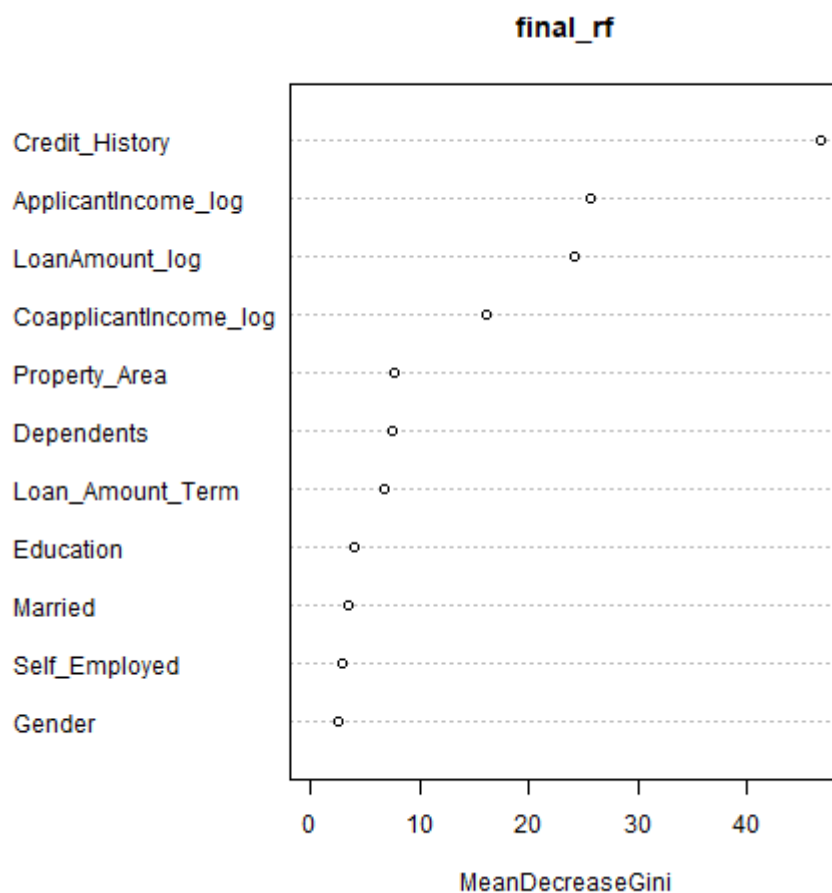
#model obtained with parameter tuning is the best fitted model as compared

Important Features

- png("FeatureImpPlot.png")
- featureImp_plt <- varImpPlot(final_rf)
- dev.off()

#Important Features are :

#Credit_History(strongest of all) , ApplicantIncome_log ,LoanAmount_log ,
CoapplicantIncome_log



4. GRADIENT BOOSTING

- `set.seed(101)`
- `xgmat <- data.matrix(trainData[-12])` #convert dataframe to matrix
- `xgmat_test <- data.matrix(testData[-12])`
- `xgb <- xgboost(data=xgmat,label = trainData$Loan_Status,nrounds =100)`

```
> xgb <- xgboost(data=xgmat,label = trainData$Loan_Status,nrounds =100)
[1] train-rmse:0.422693
[2] train-rmse:0.370246
[3] train-rmse:0.331519
[4] train-rmse:0.305919
[5] train-rmse:0.283536
[6] train-rmse:0.270067
[7] train-rmse:0.259062
[8] train-rmse:0.248201
[9] train-rmse:0.242699
[10] train-rmse:0.233021
[11] train-rmse:0.225370
[12] train-rmse:0.207365
[13] train-rmse:0.200783
[14] train-rmse:0.195978
```

- `xgb`

```
> xgb
##### xgb.Booster
raw: 383.2 Kb
call:
  xgb.train(params = params, data = dtrain, nrounds = nrounds,
    watchlist = watchlist, verbose = verbose, print_every_n = print_every_n,
    early_stopping_rounds = early_stopping_rounds, maximize = maximize,
    save_period = save_period, save_name = save_name, xgb_model = xgb_model,
    callbacks = callbacks)
params (as set within xgb.train):
  validate_parameters = "TRUE"
xgb.attributes:
  niter
callbacks:
  cb.print.evaluation(period = print_every_n)
  cb.evaluation.log()
# of features: 11
niter: 100
nfeatures : 11
evaluation_log:
  iter train_rmse
    1    0.422693
    2    0.370246
---
    99    0.013028
   100    0.012902
> |
```

- `xgb_pred <- predict(xgb,newdata = xgmat_test)`
- `xgb_pclass <- ifelse(xgb_pred<0.5,0,1)`
- `confusionMatrix(table(as.factor(testData$Loan_Status),xgb_pclass),positive = '1')`

```
Confusion Matrix and Statistics

      xgb_pclass
      0      1
0 24 25
1 13 92

              Accuracy : 0.7532
              95% CI   : (0.6774, 0.8191)
    No Information Rate : 0.7597
    P-Value [Acc > NIR] : 0.61707

              Kappa : 0.3916

  Mcnemar's Test P-Value : 0.07435

              Sensitivity : 0.7863
              Specificity : 0.6486
               Pos Pred Value : 0.8762
               Neg Pred Value : 0.4898
                Prevalence : 0.7597
                Detection Rate : 0.5974
    Detection Prevalence : 0.6818
         Balanced Accuracy : 0.7175

      'Positive' Class : 1
```

5. PREDICTIONS

#prediction

- `predictions <- predict(final_rf,af_scale)`
- `predClass <- ifelse(predictions ==1,'Y','N')`

#Final Submission score 0.7708333333333333

SUMMARY

- This is the end of the analysis, we started from data cleaning and processing, missing value imputation , then exploratory analysis and feature engineering, and finally model building and evaluation.
- The best accuracy we obtained on our validation data is 0.7792 ,Since dataset is not huge enough models performance does not vary but Random forest performed with hyperparameter tuning.
- The insights about loan approval status from the analysis is:
 - Applicants with credit history not passing guidelines mostly fails to get approved, probably because that they have a higher probability of not paying back.
 - Most of the time, applicants with high income, loaning low amount is more likely to get approved, which makes sense, those applicants are more likely to pay back their loans.
 - Having a strong co-applicant can be a plus to the probability of getting approve.